

## Lecture 06

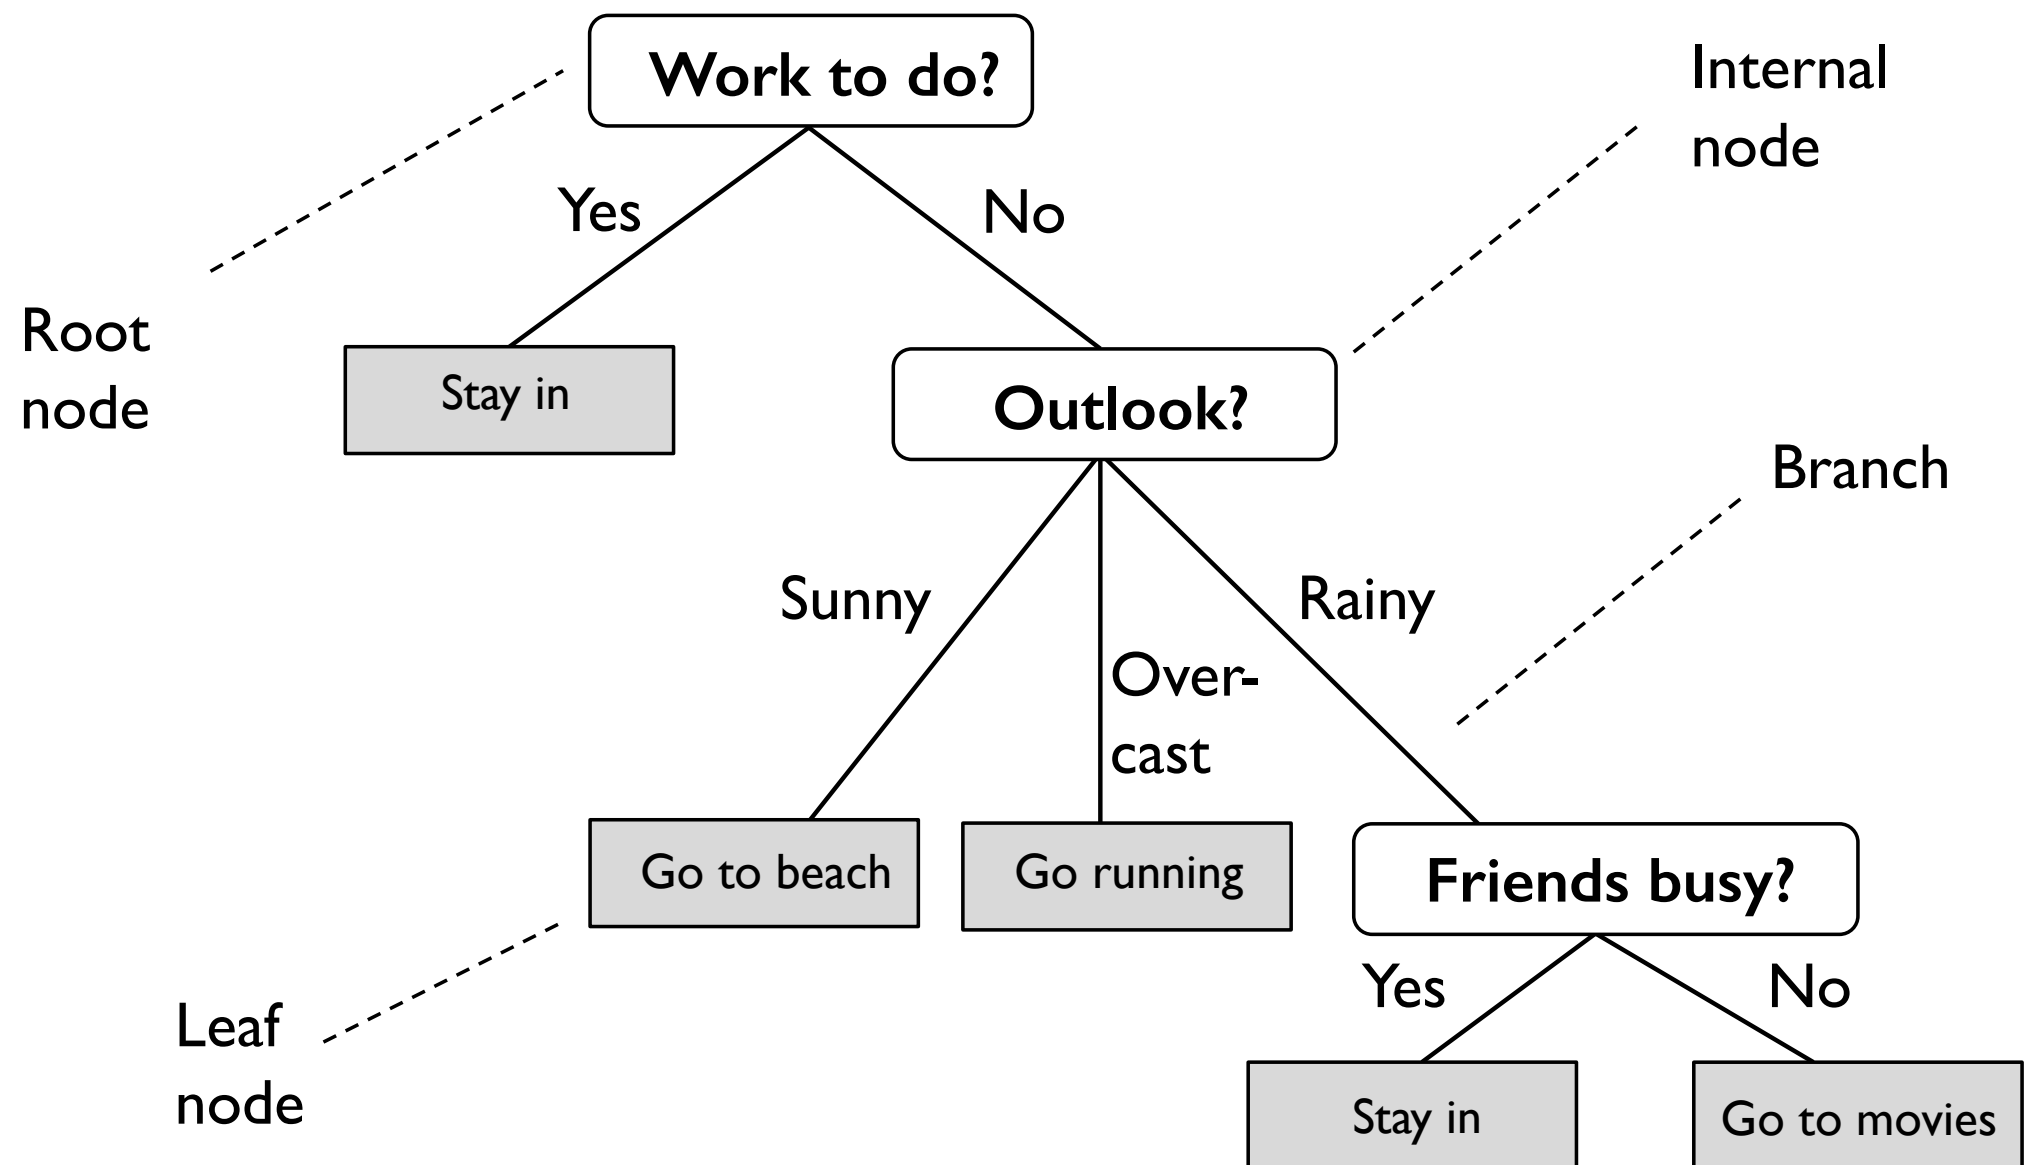
# Decision Trees

STAT 479: Machine Learning, Fall 2019

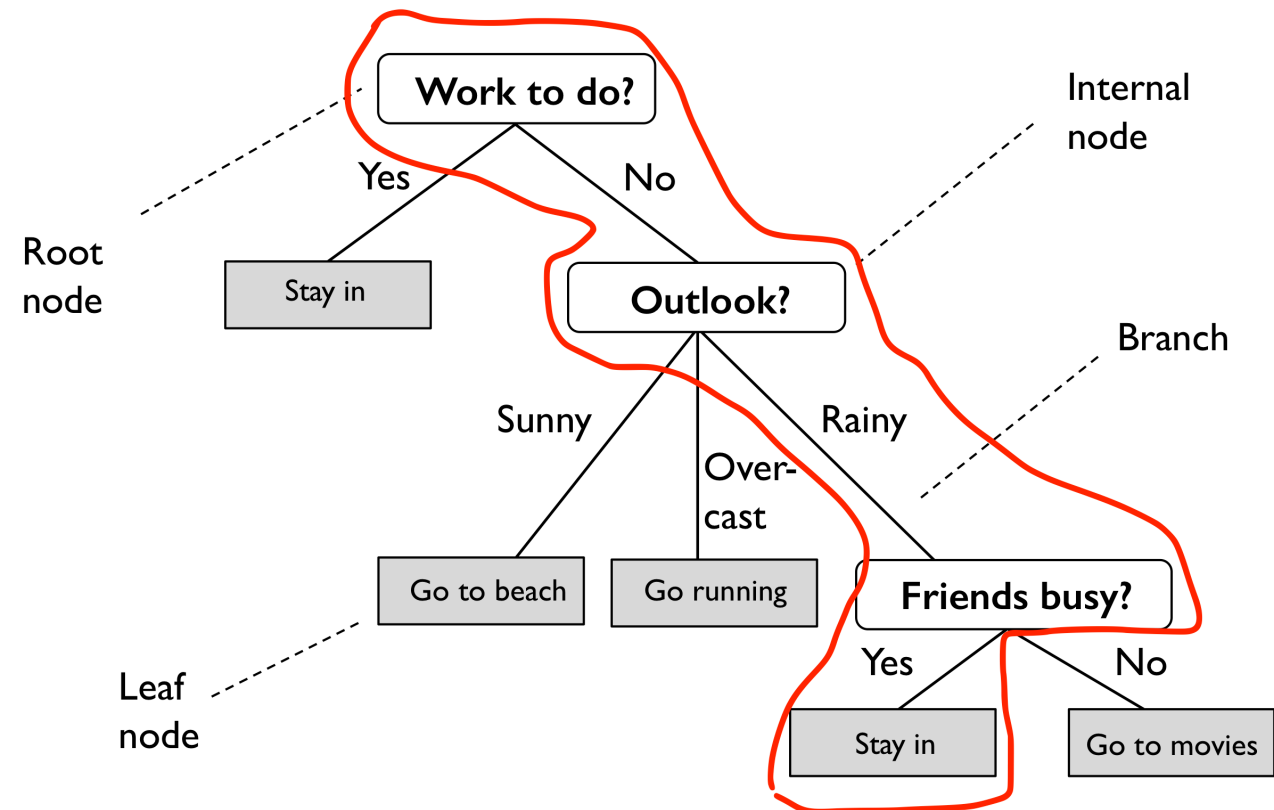
Sebastian Raschka

<http://stat.wisc.edu/~sraschka/teaching/stat479-fs2019/>

# Decision Tree Terminology



# Decision Trees as Rulesets



**IF**

\_\_\_\_\_

\_\_\_\_\_

**THEN**

\_\_\_\_\_

# Decision Trees and ML Categories

- Supervised vs. unsupervised learning algorithm
- Classification vs. regression
- Optimization method: \_\_\_\_\_
- Eager vs. lazy learning algorithm
- Batch vs. online learning algorithm
- Parametric vs. nonparametric model
- Deterministic vs. stochastic

# Recursion / Recursive Algorithms

```
1 def some_fun(x):  
2     if x == []:  
3         return 0  
4     else:  
5         return 1 + some_fun(x[1:])
```

What does this function do?

# Divide & Conquer Algorithms: Quicksort

```
1 def quicksort(array):
2     if len(array) < 2:
3         return array
4     else:
5         pivot = array[0]
6         smaller, bigger = [], []
7         for ele in array[1:]:
8             if ele <= pivot:
9                 smaller.append(ele)
10            else:
11                bigger.append(ele)
12        return quicksort(smaller) + [pivot] + quicksort(bigger)
```

# Divide & Conquer Algorithms: Quicksort

```
1 def quicksort(array):
2     if len(array) < 2:
3         return array
4     else:
5         pivot = array[0]
6         smaller, bigger = [], []
7         for ele in array[1:]:
8             if ele <= pivot:
9                 smaller.append(ele)
10            else:
11                bigger.append(ele)
12     return quicksort(smaller) + [pivot] + quicksort(bigger)
```



# Time complexity of quicksort:

$\mathcal{O}(\underline{\hspace{2cm}})$  ("on average")

```
1 def quicksort(array):
2     if len(array) < 2:
3         return array
4     else:
5         pivot = array[0]
6         smaller, bigger = [], []
7         for ele in array[1:]:
8             if ele <= pivot:
9                 smaller.append(ele)
10            else:
11                bigger.append(ele)
12        return quicksort(smaller) + [pivot] + quicksort(bigger)
```



# Array Sorting Algorithms

Algorithm	Time Complexity			Space Complexity
	Best	Average	Worst*	
<u>Quicksort</u>	$\Omega(n \log(n))$	$\Theta(n \log(n))$	$O(n^2)$	$O(\log(n))$
<u>Mergesort</u>	$\Omega(n \log(n))$	$\Theta(n \log(n))$	$O(n \log(n))$	$O(n)$
<u>Timsort</u>	$\Omega(n)$	$\Theta(n \log(n))$	$O(n \log(n))$	$O(n)$
<u>Heapsort</u>	$\Omega(n \log(n))$	$\Theta(n \log(n))$	$O(n \log(n))$	$O(1)$
<u>Bubble Sort</u>	$\Omega(n)$	$\Theta(n^2)$	$O(n^2)$	$O(1)$
<u>Insertion Sort</u>	$\Omega(n)$	$\Theta(n^2)$	$O(n^2)$	$O(1)$
<u>Selection Sort</u>	$\Omega(n^2)$	$\Theta(n^2)$	$O(n^2)$	$O(1)$
<u>Tree Sort</u>	$\Omega(n \log(n))$	$\Theta(n \log(n))$	$O(n^2)$	$O(n)$
<u>Shell Sort</u>	$\Omega(n \log(n))$	$\Theta(n(\log(n))^2)$	$O(n(\log(n))^2)$	$O(1)$
<u>Bucket Sort</u>	$\Omega(n+k)$	$\Theta(n+k)$	$O(n^2)$	$O(n)$
<u>Radix Sort</u>	$\Omega(nk)$	$\Theta(nk)$	$O(nk)$	$O(n+k)$
<u>Counting Sort</u>	$\Omega(n+k)$	$\Theta(n+k)$	$O(n+k)$	$O(k)$
<u>Cubesort</u>	$\Omega(n)$	$\Theta(n \log(n))$	$O(n \log(n))$	$O(n)$

<http://www.bigocheatsheet.com>

\* "worst" ~ inversely-sorted array

# Time Complexity ("Big-O")

Growing the tree:  $O(\dots)$

Tip: It can be shown that optimal split is on boundary between adjacent examples (similar feature value) with different class labels.

Fayyad, Usama Mohammad. "On the induction of decision trees for multiple concept learning." (1992).

# Time Complexity ("Big-O")

Querying the tree:  $\mathcal{O}(\dots)$

# Decision Tree in Pseudocode

GenerateTree( $\mathcal{D}$ ):

- if  $y = 1 \ \forall \ \langle \mathbf{x}, \mathbf{y} \rangle \in \mathcal{D}$  or  $y = 0 \ \forall \ \langle \mathbf{x}, \mathbf{y} \rangle \in \mathcal{D}$ :
  - return Tree
- else:
  - Pick best feature  $x_j$ :
    - $\mathcal{D}_0$  at Child<sub>0</sub> :  $x_j = 0 \ \forall \ \langle \mathbf{x}, \mathbf{y} \rangle \in \mathcal{D}$
    - $\mathcal{D}_1$  at Child<sub>1</sub> :  $x_j = 1 \ \forall \ \langle \mathbf{x}, \mathbf{y} \rangle \in \mathcal{D}$

return Node( $x_j$ , GenerateTree( $\mathcal{D}_0$ ), GenerateTree( $\mathcal{D}_1$ ))

# Generic Tree Growing Algorithm

- 1)** Pick the feature that, when parent node is split, results in the largest information gain
- 2)** Stop if child nodes are pure or information gain  $\leq 0$
- 3)** Go back to step 1 for each of the two child nodes

# Generic Tree Growing Algorithm

- 1) Pick the feature that, when parent node is split, results in the largest information gain
  - 2) Stop if child nodes are pure or information gain  $\leq 0$
  - 3) Go back to step 1 for each of the two child nodes
- How make predictions of features in dataset not sufficient to make child nodes pure?

# Design choices

- How to split
  - what measurement/criterion as measure of goodness
  - binary vs multi-category split
- When to stop
  - if leaf nodes contain only examples of the same class
  - feature values are all the same for all examples
  - statistical significance test

# ID3 -- Iterative Dichotomizer 3

- one of the earlier/earliest decision tree algorithms
- Quinlan, J. R. 1986. Induction of Decision Trees. Mach. Learn. 1, 1 (Mar. 1986), 81-106.
- cannot handle numeric features
- no pruning, prone to overfitting
- short and wide trees (compared to CART)
- maximizing information gain/minimizing entropy
- discrete features, binary and multi-category features



# C4.5

- continuous and discrete features
- Ross Quinlan 1993, Quinlan, J. R. (1993). C4.5: Programming for machine learning. *Morgan Kauffmann*, 38, 48.
- continuous is very expensive, because must consider all possible ranges
- handles missing attributes (ignores them in gain compute)
- post-pruning (bottom-up pruning)
- Gain Ratio

# CART

- Breiman, L. (1984). *Classification and regression trees*. Belmont, Calif: Wadsworth International Group.
- continuous and discrete features
- strictly binary splits (taller trees than ID3, C4.5)
- binary splits can generate better trees than C4.5, but tend to be larger and harder to interpret; k-attributes has a ways to create a binary partitioning
- variance reduction in regression trees
- Gini impurity, twoing criteria in classification trees
- cost complexity pruning

# Others

- CHAID (CHi-squared Automatic Interaction Detector); Kass, G. V. (1980). "An exploratory technique for investigating large quantities of categorical data". *Applied Statistics*. 29 (2): 119–127.
- MARS (Multivariate adaptive regression splines); Friedman, J. H. (1991). "Multivariate Adaptive Regression Splines". *The Annals of Statistics*. 19: 1
- C5.0 (patented)
- ...

# 2019 UW Data Challenge Instructions

Organized by the Undergraduate Statistics Club

This event provides an opportunity for students to gain practice in working with real data and communicating the analysis results clearly and concisely. Positive outcomes of this data challenge will be rewarded based on an oral presentation and a short two-page executive report.

## **Timeline for Data Challenge:**

Slides, Report, and Code Submission Deadline:

Tuesday October 8<sup>th</sup> at 5:00 pm

Final presentation:

Tuesday October 8<sup>th</sup> at 6:00 pm

Room: Service Memorial Institute (SMI) 133

Please arrive 15 minutes early with your team to prepare.

---

## **Dataset:**

For this data challenge, we are using the data from an open data competition on Kaggle.com. The dataset is about predicting a movie's world-wide box office revenue based on metadata on over 7,000 past films (e.g. budget, genre, runtime, cast, etc).

## [DataSci] Machine Learning Series

Scheduled: Thursday Oct 10, 6:30 PM

Location: Genetics-Biotech Center 1441



UW DATA SCIENCE CLUB

Machine Learning Series :

### Soft-Biometric Attributes Prediction from Face Images with PyTorch

Sebastian Raschka

Thursday Oct 10, 6:30pm  
eGenetics-Biotech 1441

## Predicting soft- biometric attributes from face images with deep learning and PyTorch

Soft-biometric characteristics include a person's age, gender, race, and health status. As many Deep Learning-centric applications are developed in recent years, the automatic extraction of soft biometric attributes can happen without the user's agreement, thereby raising several privacy concerns. This talk will introduce how to extract soft-biometric attributes from facial images, as well as how to conceal soft-biometric information for enhancing privacy.

**Don't worry if you do not have programming experience with Python!** Dr. Rashka will also give a tutorial introducing PyTorch and how we can use it to train a simple gender classifier and ordinal regression model for estimating the apparent age from face images.

Best,


Lareina Liu  
UW Data Science Club

**.Data**


# Finding a Decision Rule

$x_1$	$x_2$	$x_3$	$y$
6 cm	8 cm	9 cm	1
4 cm	11 cm	2 cm	0
6 cm	12 cm	4 cm	0
10 cm	9 cm	3 cm	1
5 cm	7 cm	8 cm	0
8 cm	9 cm	3 cm	1
3 cm	11 cm	5 cm	0

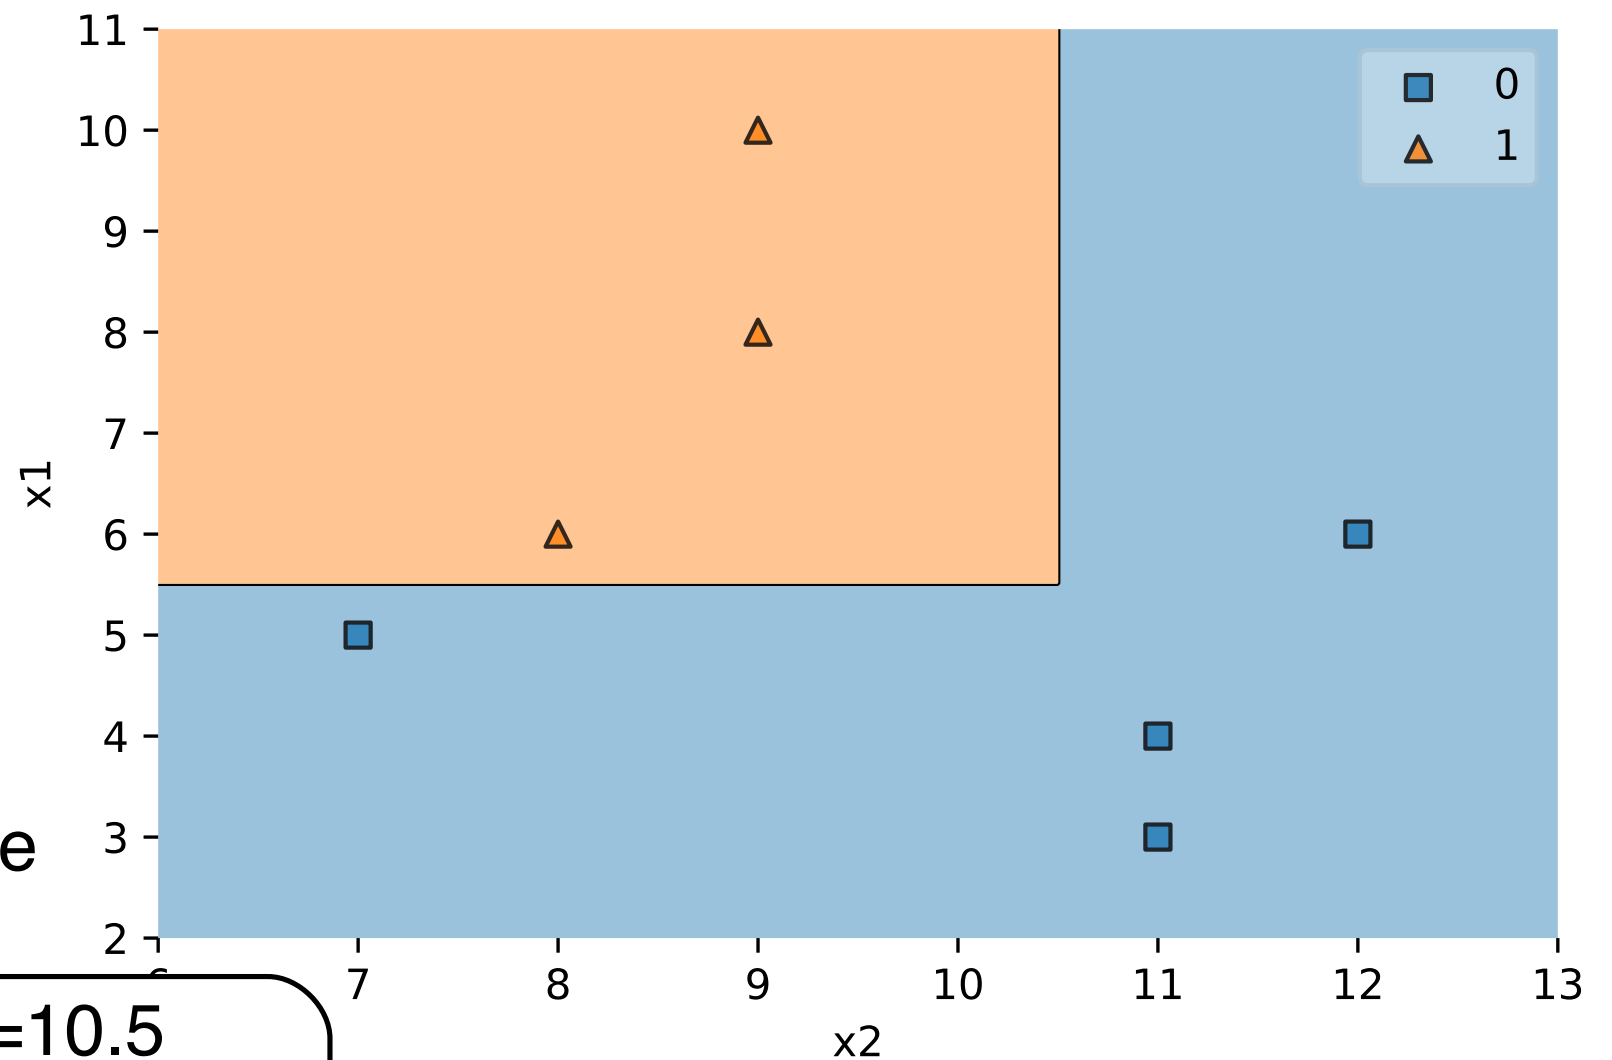
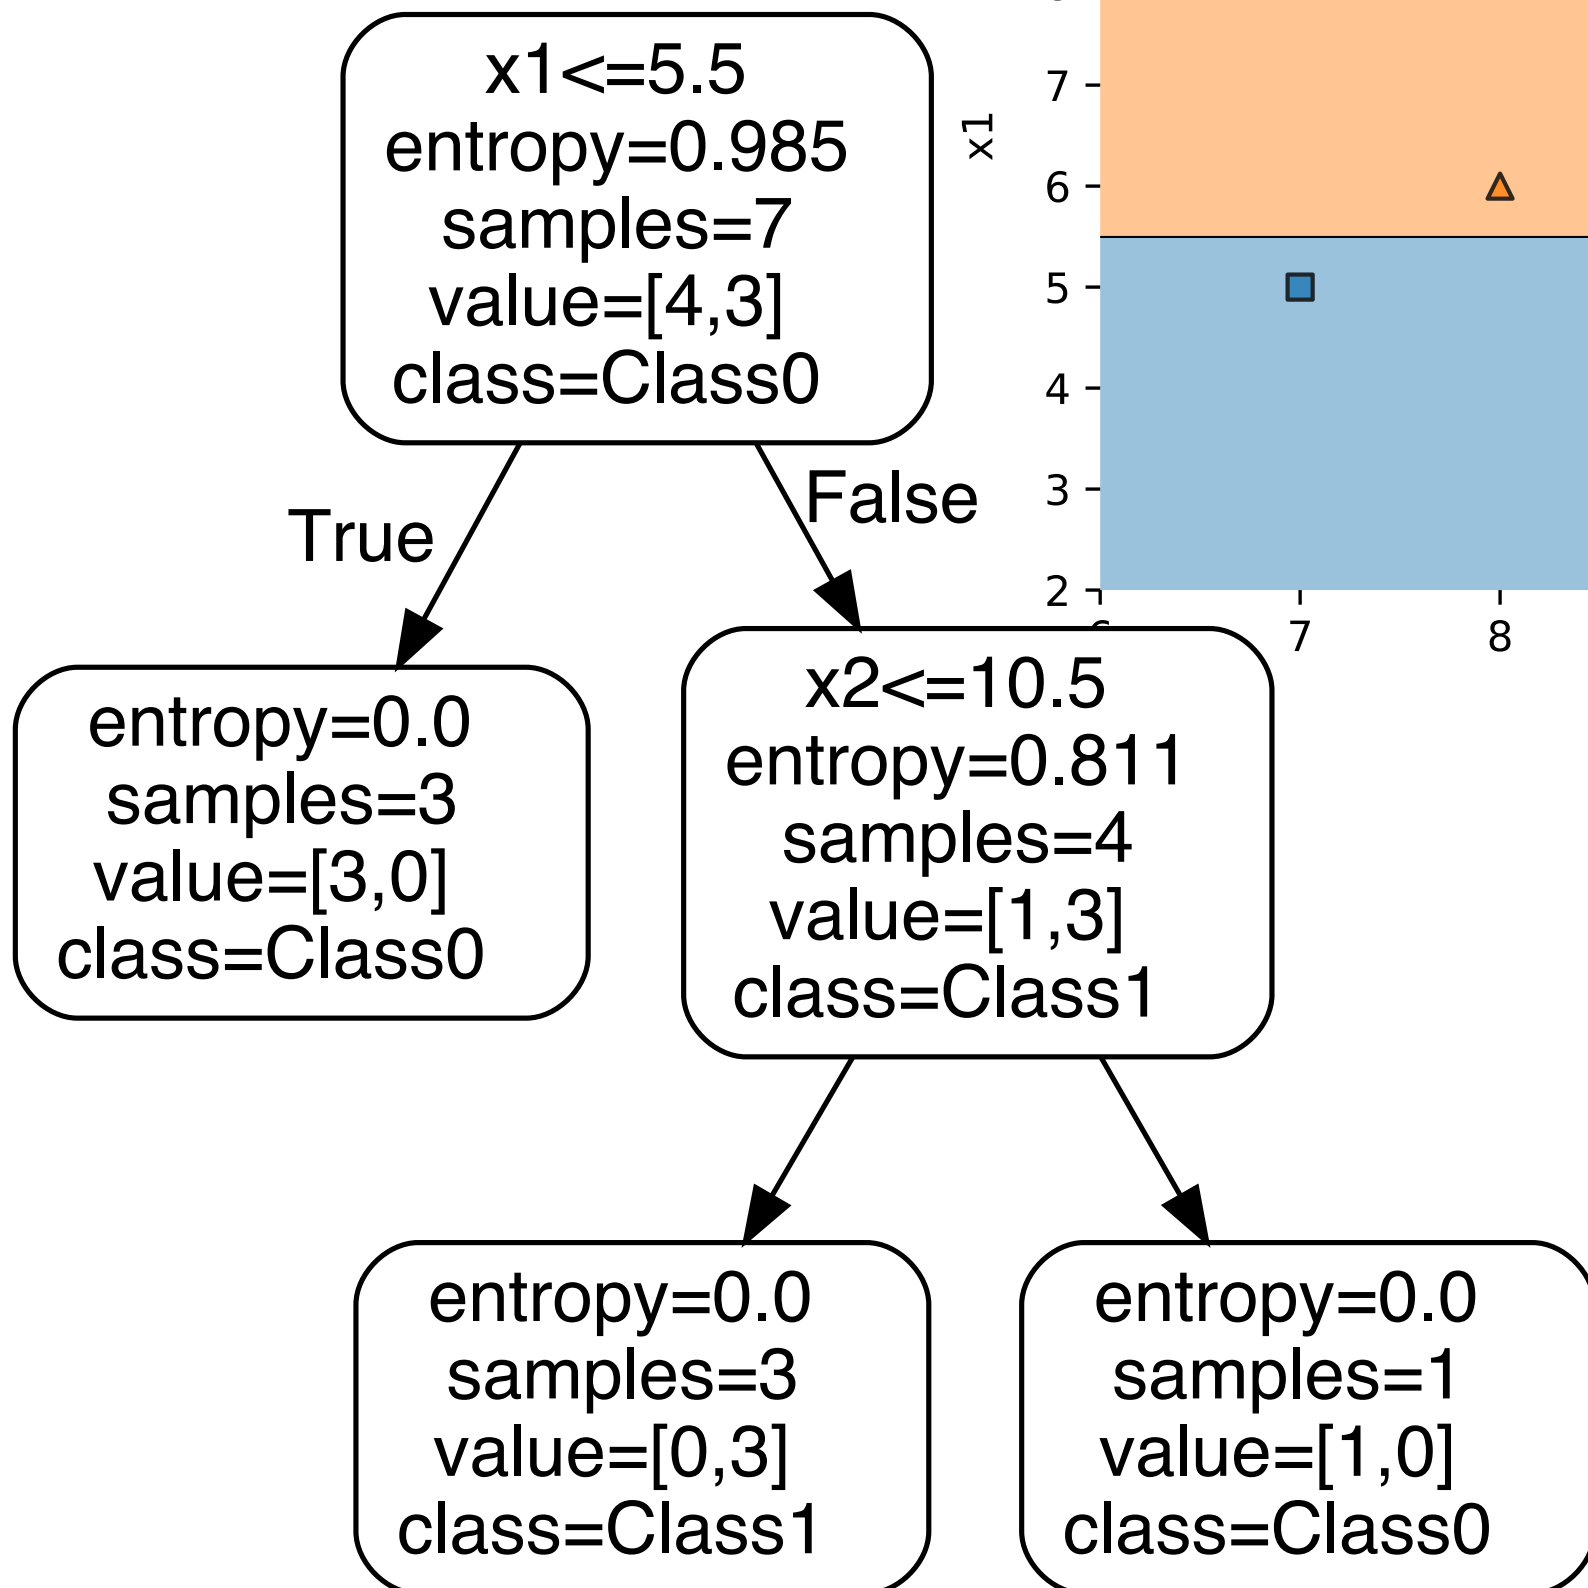
# Drawing a Decision Boundary




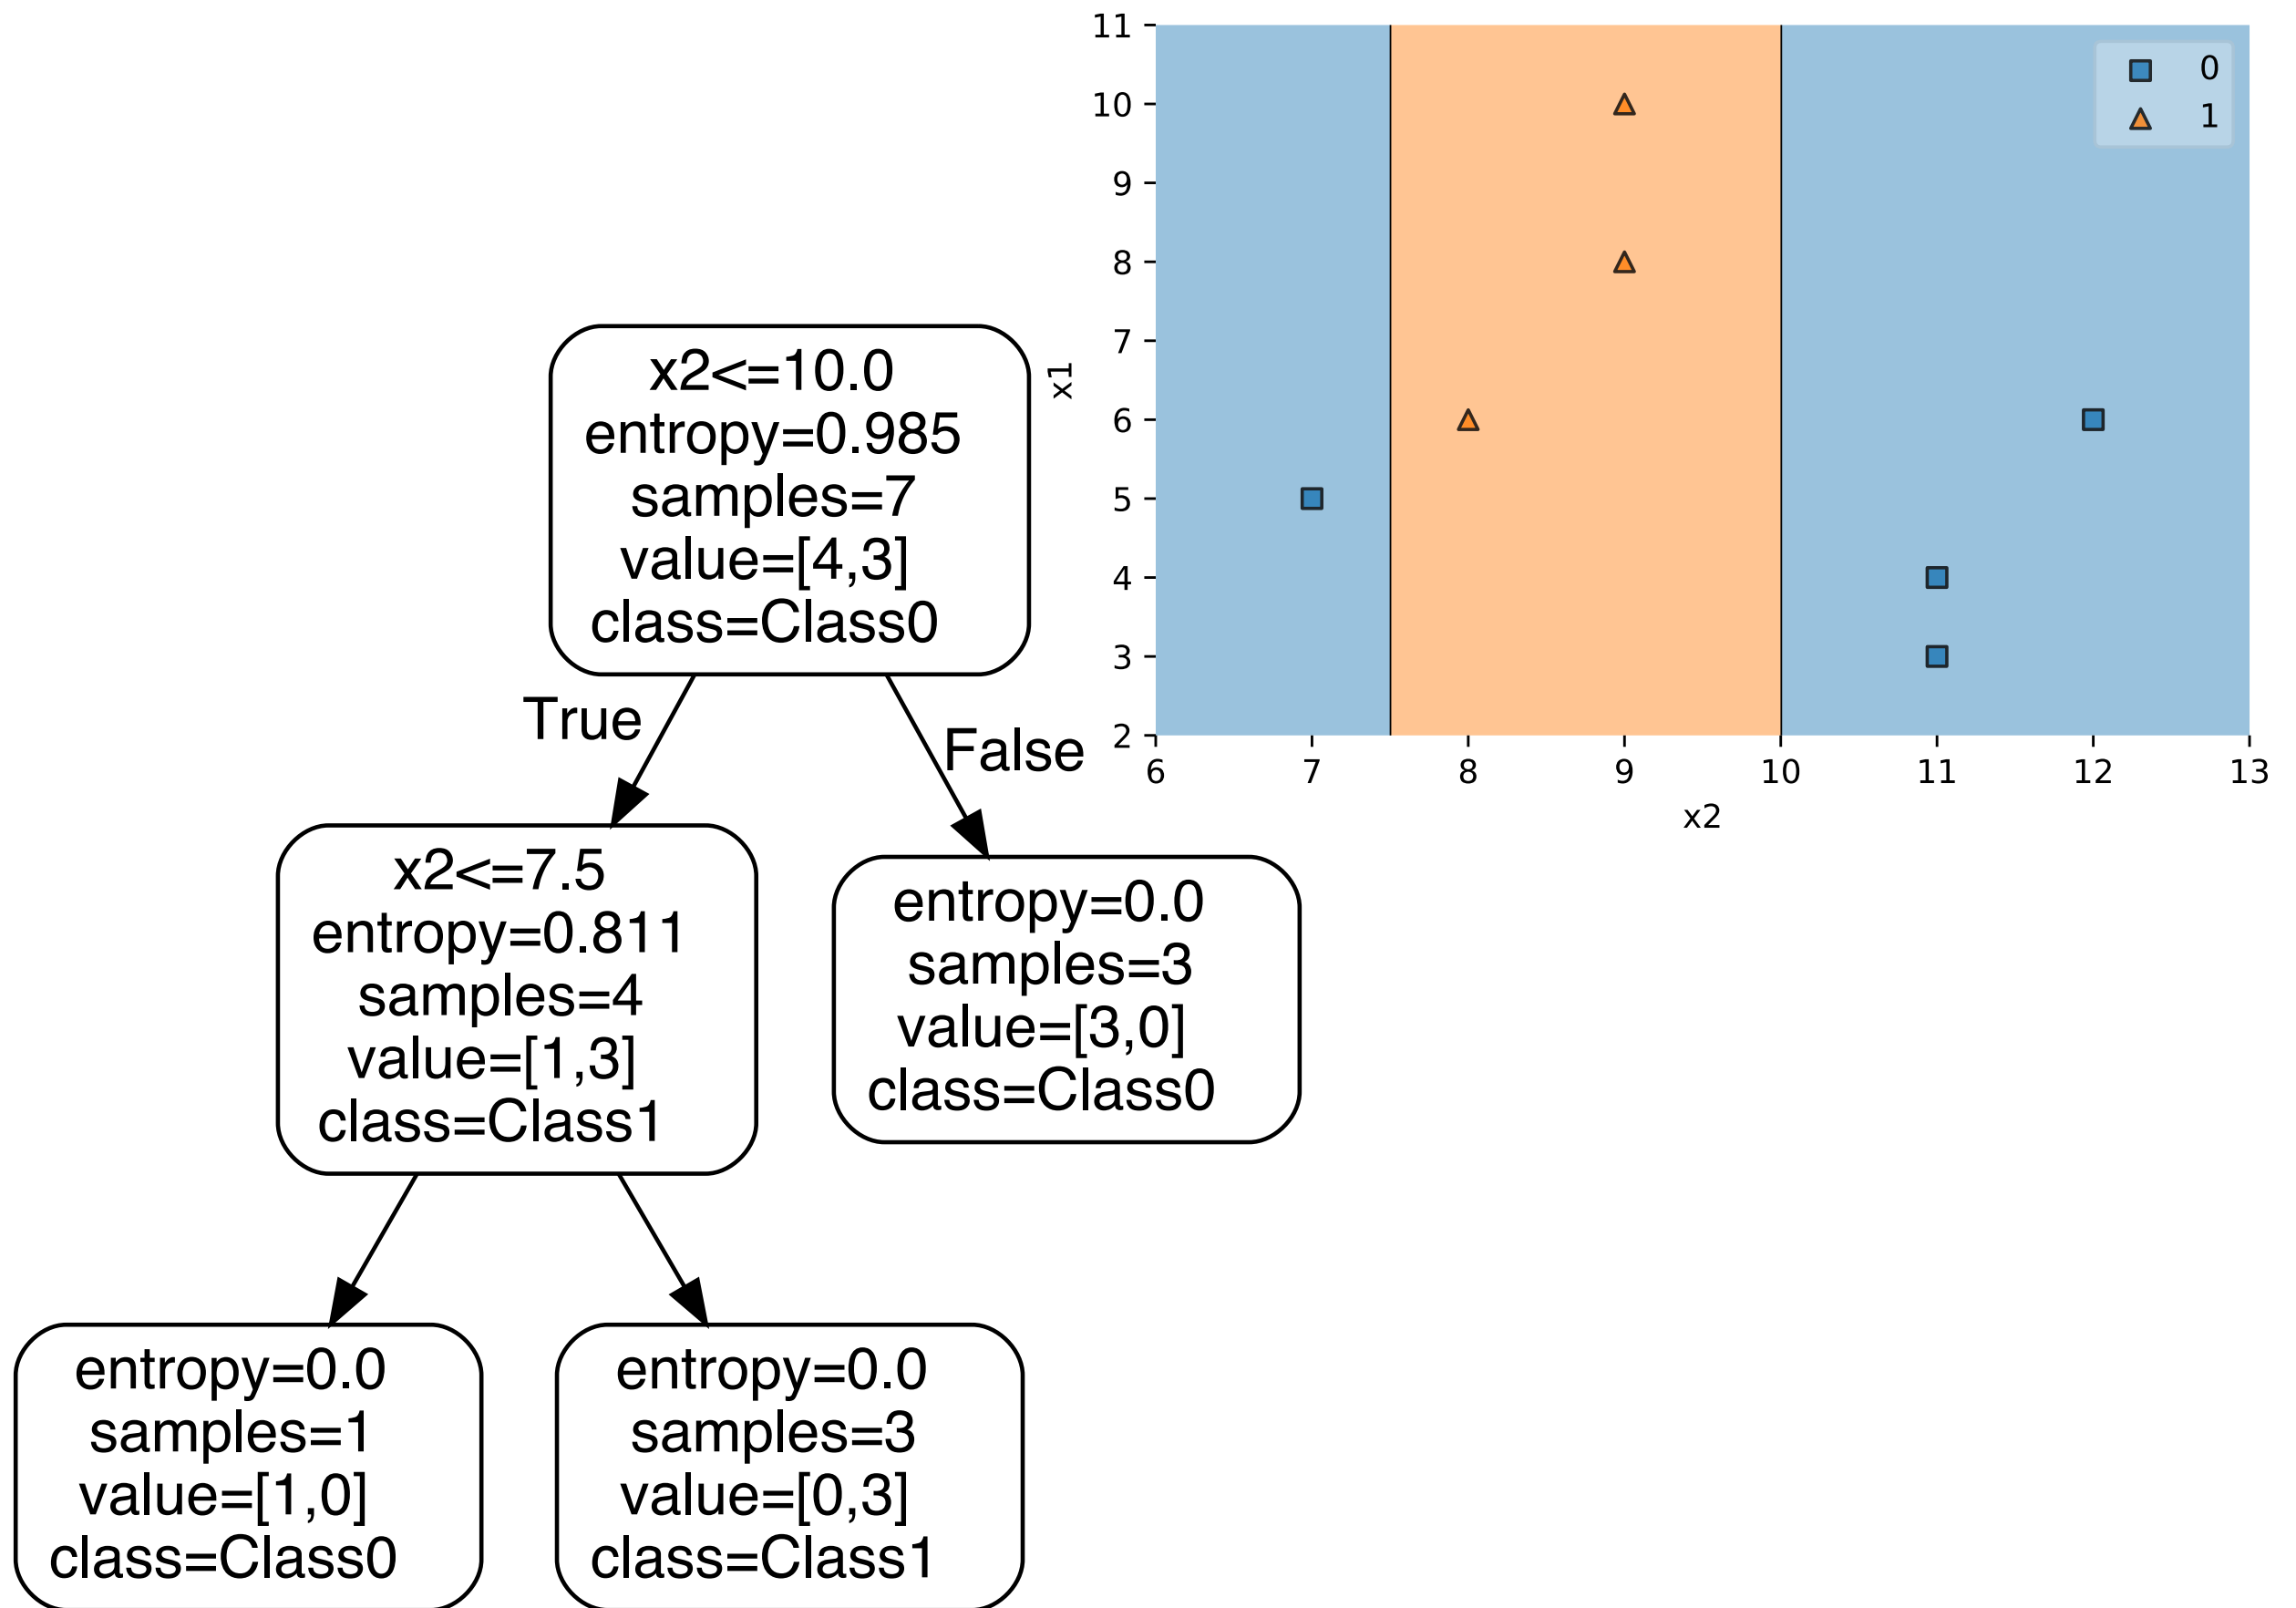
$x_1$	$x_2$	$x_3$	$y$
6 cm	8 cm	9 cm	1
4 cm	11 cm	2 cm	0
6 cm	12 cm	4 cm	0
10 cm	9 cm	3 cm	1
5 cm	7 cm	8 cm	0
8 cm	9 cm	3 cm	1
3 cm	11 cm	5 cm	0











# The Splitting Criterion

# Information Gain

$$GAIN(\mathcal{D}, x_j) = H(\mathcal{D}) - \sum_{v \in Values(x_j)} \frac{|\mathcal{D}_v|}{|\mathcal{D}|} H(\mathcal{D}_v)$$

# Shannon Entropy

Refer to lecture notes

# Entropy

$$H = - \sum_i p(i | x_j) \log_2(p(i | x_j))$$

# Gini Impurity

$$Gini = 1 - \sum_i (p(i | x_j)^2)$$




# Misclassification Error

$$ERR = \frac{1}{n} \sum_{i=1}^n L(\hat{y}^{[i]}, y^{[i]}),$$

$$L(\hat{y}, y) = \begin{cases} 0 & \text{if } \hat{y} = y, \\ 1 & \text{otherwise.} \end{cases}$$

# Misclassification Error

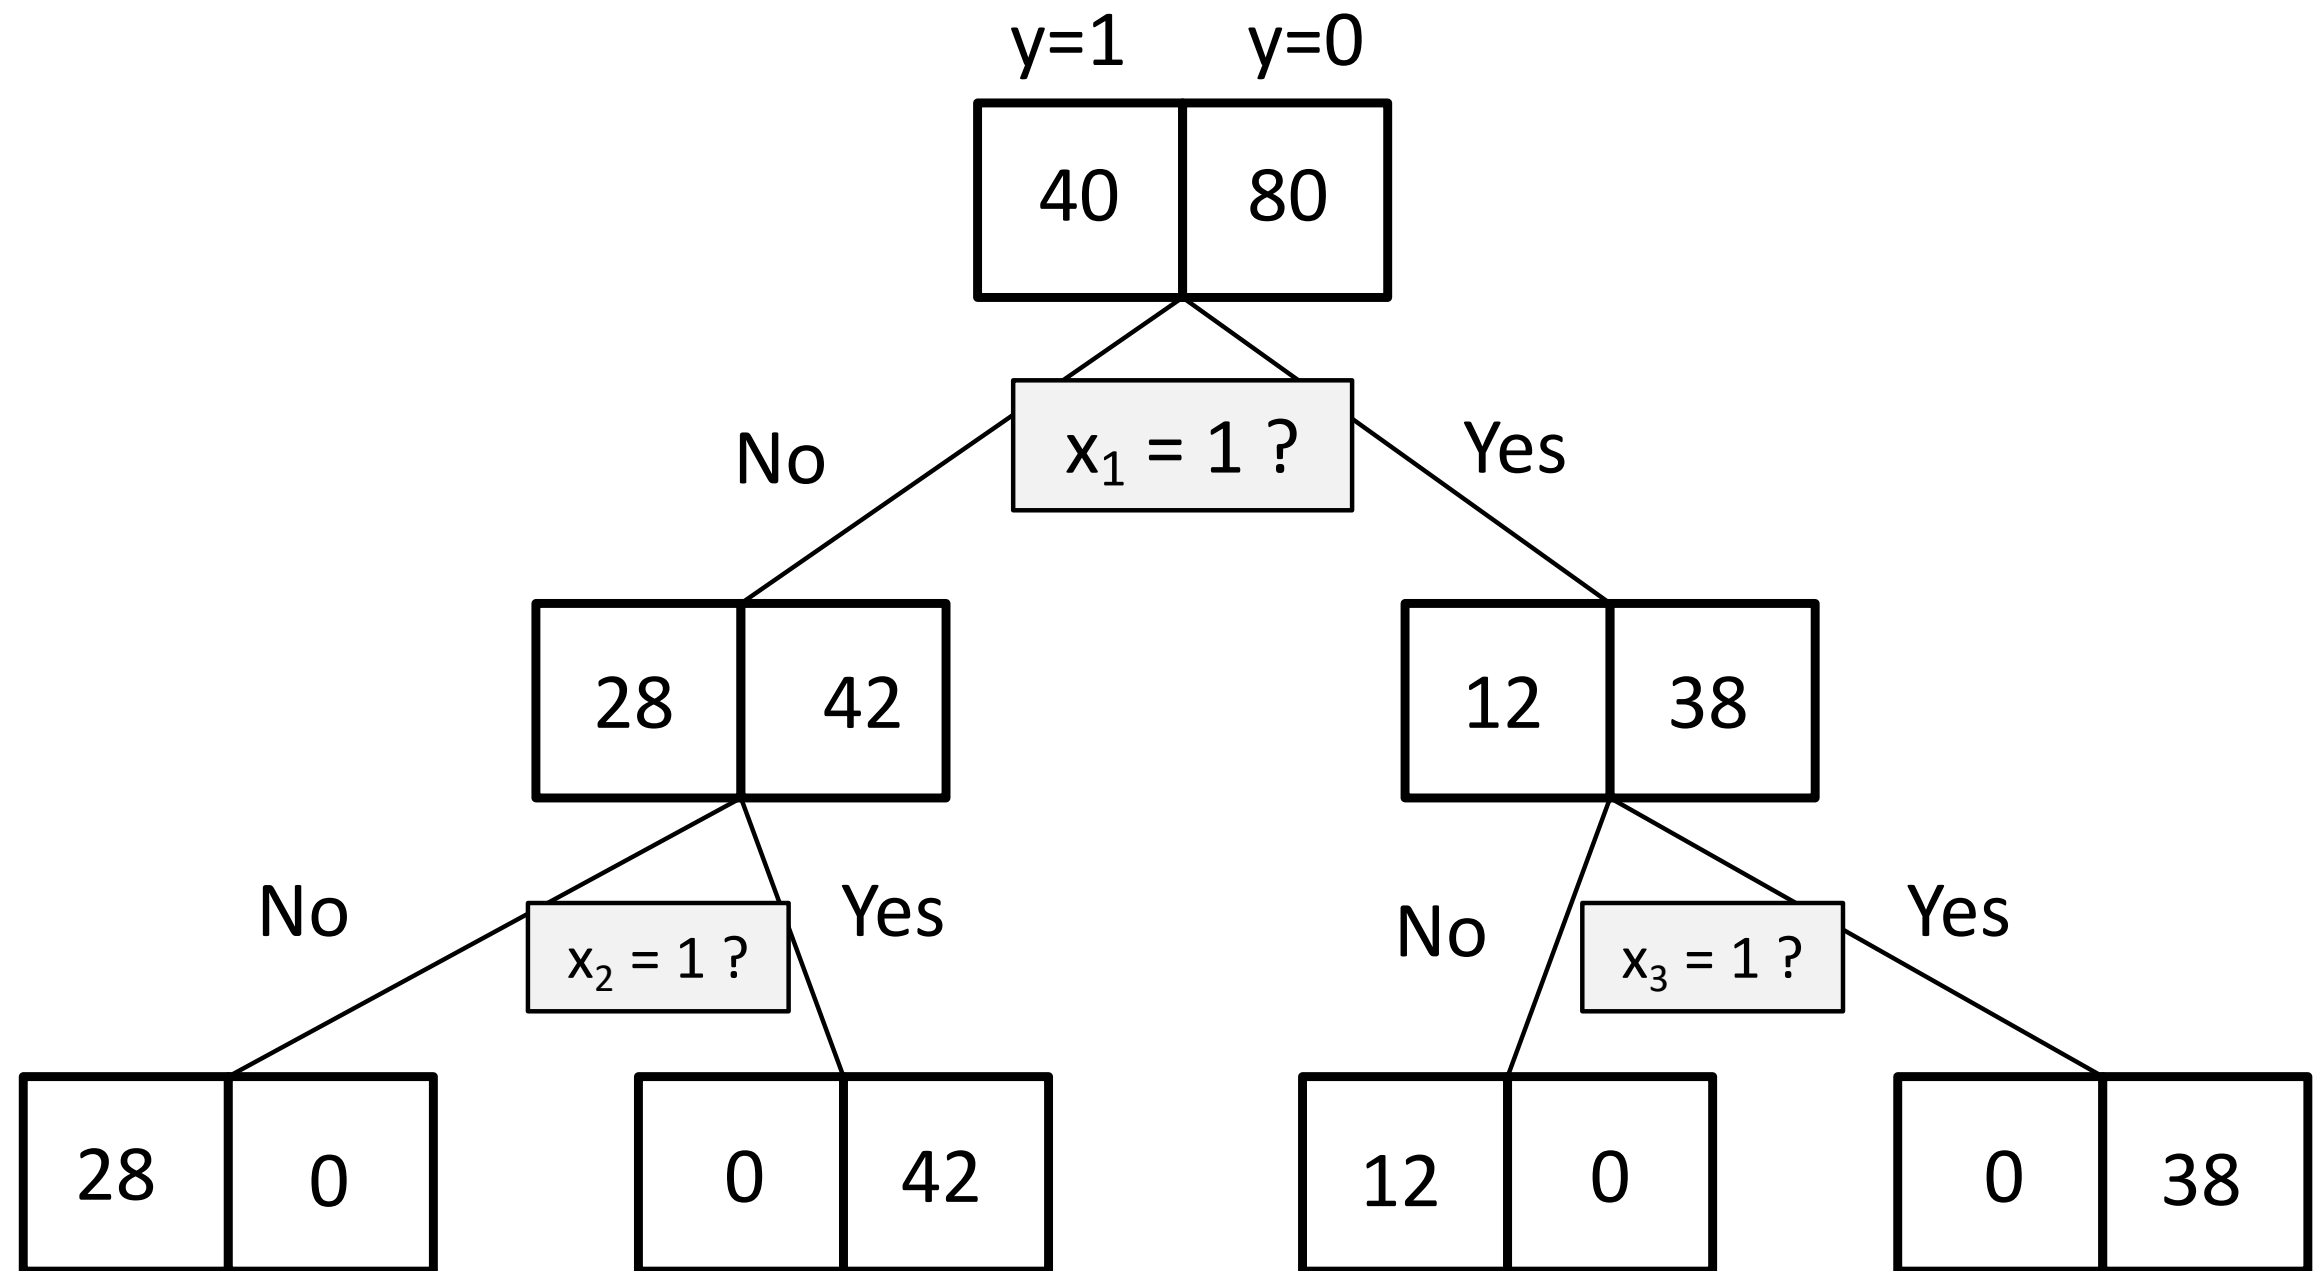
$$ERR = 1 - \max_i(p(i | x_j))$$

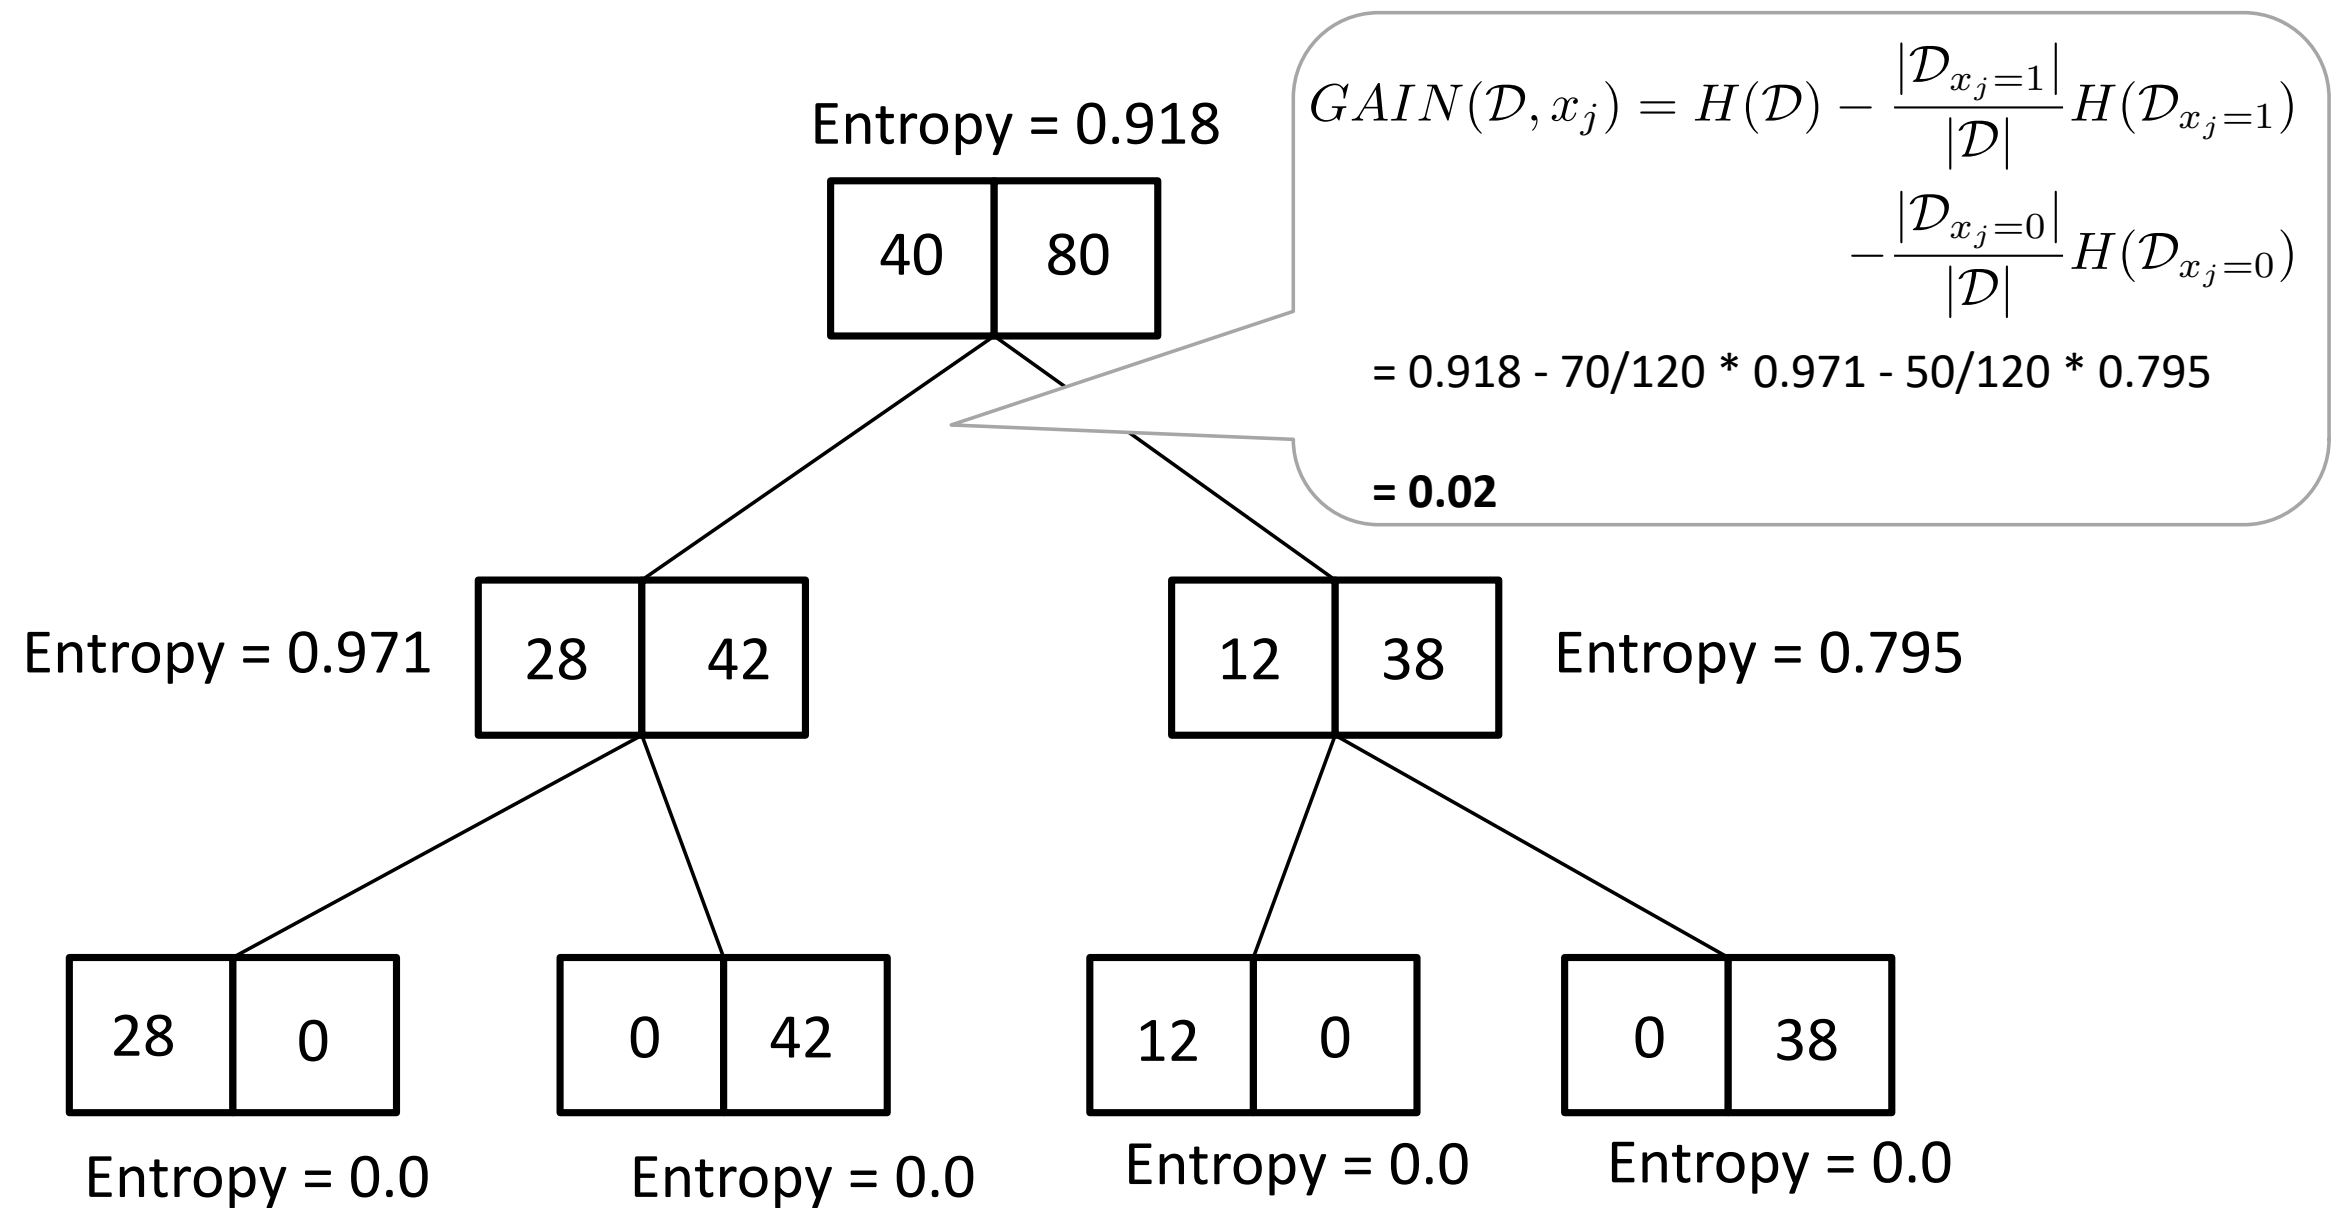


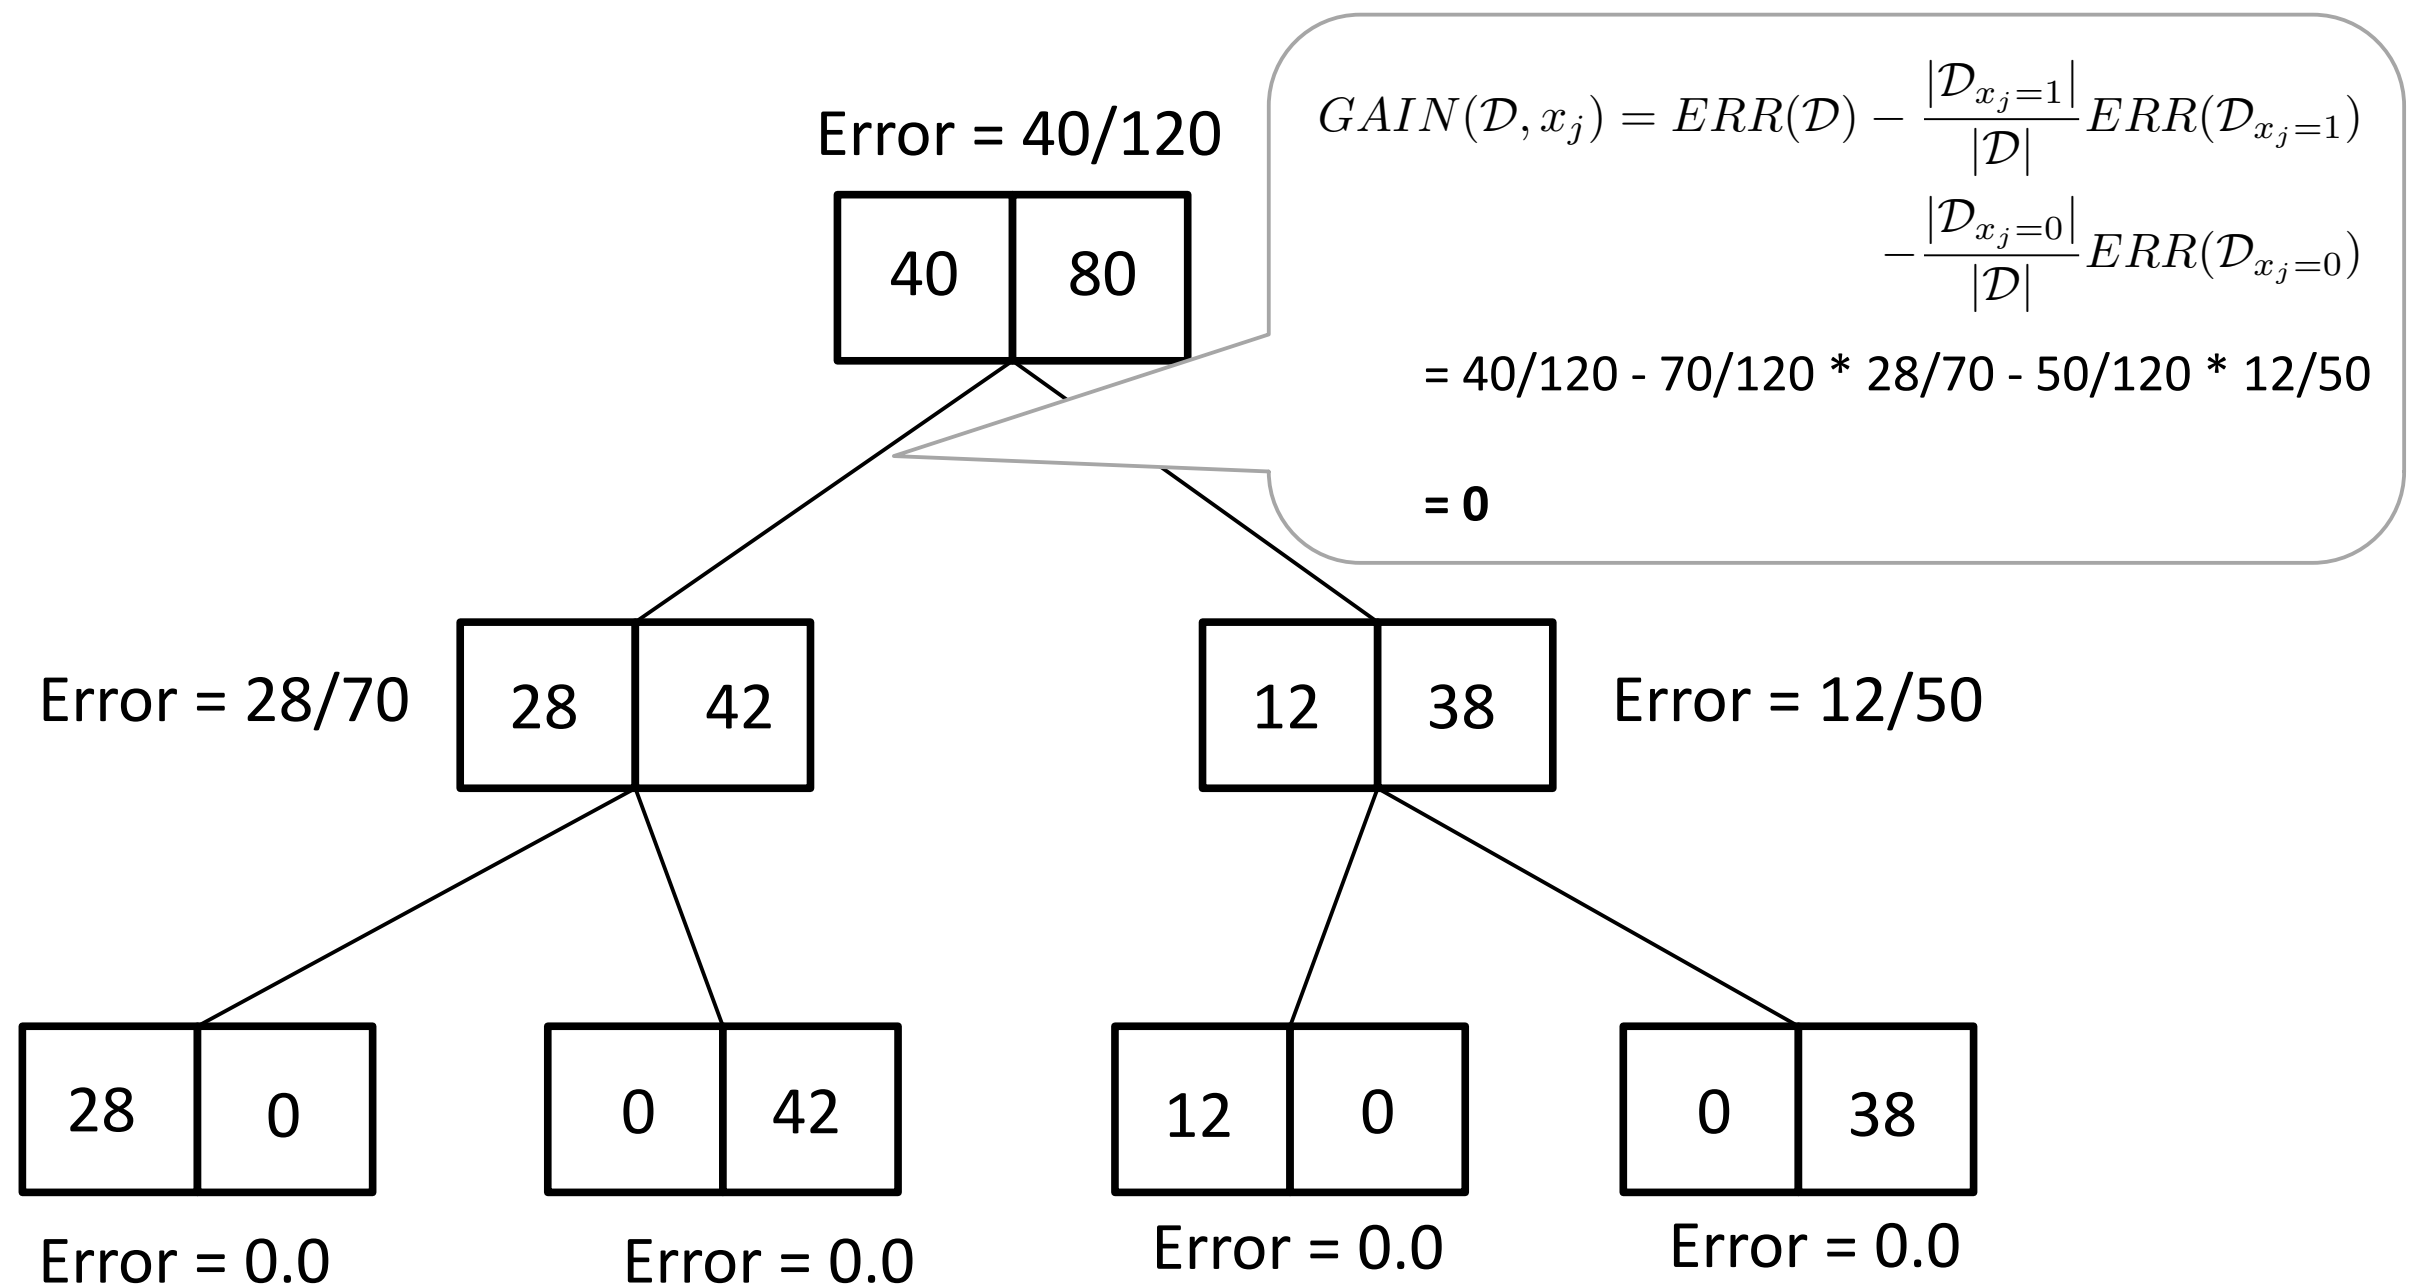
# **Why Growing Decision Trees via Entropy instead of Misclassification Error?**

# Why Growing Decision Trees via Entropy instead of Misclassification Error?


$$GAIN(\mathcal{D}, x_j) = I(\mathcal{D}) - \sum_{v \in Values(x_j)} \frac{|\mathcal{D}_v|}{|\mathcal{D}|} I(\mathcal{D}_v)$$











# Gain Ratio

$$\textit{GainRatio}(\mathcal{D}, x_j) = \frac{\textit{Gain}(\mathcal{D}, x_j)}{\textit{SplitInfo}(\mathcal{D}, x_j)}$$


Quinlan 1986

where the split information measures the entropy of the feature:

$$\textit{SplitInfo}(\mathcal{D}, x_j) = - \sum_{v \in x_j} \frac{|\mathcal{D}_v|}{|\mathcal{D}|} \log_2 \frac{|\mathcal{D}_v|}{|\mathcal{D}|}$$


Penalizes splitting categorical attributes with many values (e.g., think date column, or really bad: row ID) via the split information

# Shortcomings



## How would the decision tree split look like?

# Overfitting




# Pre-Pruning

- Set a depth cut-off (maximum tree depth) *a priori*
- Cost-complexity pruning: , where is an impurity measure, is a tuning parameter, and is the total number of nodes.
- Stop growing if split is not statistically significant (e.g.,  $\chi^2$  test)
- Set a minimum number of data points for each node

# Post-Pruning

- Grow full tree first, then remove nodes, in C4.5
- Reduced-error pruning, remove nodes via validation set eval. (problematic for limited data)
- Can also convert trees to rules first and then prune the rules

# Post-Pruning



# Regression Trees



# Decision Tree Summary: Pros and Cons

- (+) Easy to interpret and communicate
- (+) Can represent "complete" hypothesis space
- (-) Easy to overfit
- (-) Elaborate pruning required
- (-) Expensive to just fit a "diagonal line"
- (-) Output range is bounded (dep. on training examples) in regression trees

# Reading Assignments

Python Machine Learning, 2nd Ed., Ch03 pg. 88-104

# Demo

06-trees\_demo.ipynb